

TACC Technical Report IMP-12

K-means clustering in the Integrative Model

Victor Eijkhout*

January 16, 2017

This technical report is a preprint of a paper intended for publication in a journal or proceedings. Since changes may be made before publication, this preprint is made available with the understanding that anyone wanting to cite or reproduce it ascertains that no published version in journal or proceedings exists.

Permission to copy this report is granted for electronic viewing and single-copy printing. Permissible uses are research and browsing. Specifically prohibited are *sales* of any copy, whether electronic or hardcopy, for any purpose. Also prohibited is copying, excerpting or extensive quoting of any report in another work without the written permission of one of the report's authors.

The University of Texas at Austin and the Texas Advanced Computing Center make no warranty, express or implied, nor assume any legal liability or responsibility for the accuracy, completeness, or usefulness of any information, apparatus, product, or process disclosed.

* eijkhout@tacc.utexas.edu, Texas Advanced Computing Center, The University of Texas at Austin

Abstract

We consider the k -means clustering algorithm.

The following IMP reports are available or under construction:

- IMP-00** The IMP Elevator Pitch
- IMP-01** IMP Distribution Theory
- IMP-02** The deep theory of the Integrative Model
- IMP-03** The type system of the Integrative Model
- IMP-04** Task execution in the Integrative Model
- IMP-05** Processors in the Integrative Model
- IMP-06** Definition of a ‘communication avoiding’ compiler in the Integrative Model (under construction)
- IMP-07** Associative messaging in the Integrative Model (under construction)
- IMP-08** Resilience in the Integrative Model (under construction)
- IMP-09** Tree codes in the Integrative Model
- IMP-10** Thoughts on models for parallelism
- IMP-11** A gentle introduction to the Integrative Model for Parallelism
- IMP-12** K-means clustering in the Integrative Model
- IMP-13** Sparse Operations in the Integrative Model for Parallelism
- IMP-14** 1.5D All-pairs Methods in the Integrative Model for Parallelism (under construction)
- IMP-15** Collectives in the Integrative Model for Parallelism
- IMP-16** Processor-local code (under construction)
- IMP-17** The CG method in the Integrative Model for Parallelism (under construction)
- IMP-18** A tutorial introduction to IMP software (under construction)
- IMP-19** Report on NSF EAGER 1451204.
- IMP-20** A mathematical formalization of data parallel operations
- IMP-21** Adaptive mesh refinement (under construction)
- IMP-22** Implementing LULESH in IMP (under construction)
- IMP-23** Distributed computing theory in IMP (under construction)
- IMP-24** IMP as a vehicle for software/hardware co-design, with John McCalpin (under construction)
- IMP-25** Dense linear algebra in IMP (under construction)

1 Introduction

1.1 Standard algorithm

Gleaned from HPC challenge [1].

The abstract algorithm iterates the following calculation:

```

for point i<N:
  for cluster k<C:
    compute distance point i to center k
  kmin is minimum distance
  cluster_i = kmin
for cluster k<C:
  let G_k the group of all i that have k as cluster number
  center_k is sum coordinates over cluster_k / |G_k|

```

Let N be the number of points and K the number of clusters, then the data structures are:

- $X(N)$ are coordinates
- $C(K)$ are cluster center coordinates,
- $d(N, K)$ are distances to cluster centers
- $g(N)$ is the group assignment.

The algorithm then consist of

- Distance computation, which is a outer product

$$d(i, k) \leftarrow f(X(i), C(k))$$

and therefore $N \times K$ way parallel.

- Group assignment, parallel over the points:

$$g(i) \leftarrow \underset{k}{\operatorname{argmin}} g(d(i, k))$$

with a reduction over the K dimension.

- Cluster formation, parallel over the clusters:

$$C(k) \leftarrow \underset{i}{\operatorname{reduce}} h(d(i, k), g(i))$$

with a reduction over the N dimension.

Here f, g, h are scalar functions with obvious but irrelevant definitions.

Let u be a distribution of N over P and let v be a distribution of K over P . Then the general calculation becomes:




Figure 1: Illustration of the distributions in a kmeans algorithm

$$\begin{array}{ll}
 d(u, v) = X(u) \cdot C(v) & \text{parallel pointwise mult from redundant } C \\
 g(u) = \operatorname{argmin}_2 d(u, \star) & \text{parallel reduction} \\
 G(u, v) = (i, k) \mapsto \delta(g(i) = k) & \text{definition of group mask} \\
 C(v) = \operatorname{reduce}_i(X(i), G(i, k)) & \text{new center from masked computation}
 \end{array}$$

If the number of clusters is small compared to the number of points or the number of processors, the easiest parallelization strategy is to replicate the k dimension and only divide the N dimension.

$$\begin{array}{ll}
 d(u, \star) = X(u) \cdot C(\star) & d(u, \star) = X(u) \cdot C(\star) \\
 g(u) = \operatorname{argmin}_2 d(u, \star) & g(u, \star) = \text{maskarray } \operatorname{argmin}_2 d(u, \star) \\
 \forall_{k < K}: G_k = \{i: g(i) = k\} & \\
 \forall_{k < K}: \{P\}: c_k = M(G_k: X(u)) & c_k(\star) = M(g(u, \star) \otimes X(u, \star))
 \end{array}$$

1.2 Gonzalez algorithm

See [2].

```

centers <- points
while #centers too high:
  partition centers
  compute k centers in each partition
  centers <- union of all the partition centers

```

See [3] for analysis and performance of Gonzalez and other parallel methods. They use MapReduce.

2 Implementation

We have a global size N , space dimension d , and a number of centers k . The centers are stored redundantly:

```
%% template_kmeans.cxx
IMP_distribution
  *krelicated = new IMP_replicated_distribution(decomp,dim,ncluster);
IMP_object
  *centers = new IMP_object( krelicated );
```

Setting the initial centers is data parallel:

```
set_initial_centers( centers );
```

On the other hand, the coordinates are a properly distributed data set of size $N \times d$, where d is the space dimension, typically 2 or 3:

```
%% unittest_kmeans.cxx
mpi_distribution
  *twoblocked = new mpi_block_distribution
    (decomp,dim,-1,globalsize);
mpi_object
  *coordinates = new mpi_object( twoblocked );
```

2.1 Origin kernel

We generate the coordinates in an origin kernel:

```
%% unittest_kmeans.cxx
mpi_kernel
  *set_random_coordinates = new mpi_kernel( coordinates );
set_random_coordinates->localexecutefn = &generate_random_coordinates;
```

2.2 Distance to clusters

Calculate the distances to each of the k coordinates, given an $N \times dk$ array is an outer product calculation. Here we distribute the cross dimension redundantly:

```
%% unittest_kmeans.cxx
mpi_distribution
  *kblocked = new mpi_block_distribution
    (decomp,ncluster,-1,globalsize);
mpi_object
  *distances = new mpi_object( kblocked );
```

The distance calculation has two input object, both with $\beta = \gamma$:

```
%% unittest_kmeans.cxx
calculate_distances = new mpi_kernel( coordinates,distances );
calculate_distances->last_dependency()->set_explicit_beta_distribution(coordinates);
calculate_distances->add_in_object( centers);
```

```
calculate_distances->last_dependency()->set_explicit_beta_distribution( centers );
calculate_distances->set_localexecutefn( &distance_calculation);
```

We summarize this as a derived kernel class:

```
%% unittest_kmeans.cxx
printf("the outer product kernel still uses the context. Wrong!\n");
calculate_distances = new mpi_outerproduct_kernel
( coordinates,distances,centers,&distance_calculation );

%% mpi_ops.h
mpi_outerproduct_kernel( object *in,object *out,
    object *replicated,
    void (*f)(int, processor_coordinate*, std::vector<object*>*, object*, double*)
)
: mpi_kernel(in,out),entity(entity_cookie::KERNEL) {
    last_dependency()->set_explicit_beta_distribution(in);
    add_in_object(replicated);
    last_dependency()->set_explicit_beta_distribution(replicated);

    set_name(fmt::format("outer product{}",&get_out_object()->get_object_number()));
    set_localexecutefn(f);
};
```

The definition of the outer product kernel (and other collectives) is further discussed in [?].

2.3 Find nearest center

With the distances, we can compute the nearest center, an array of length N , in a local operation. The nearest center is an integer value, but we store it as real.

```
%% unittest_kmeans.cxx
mpi_distribution
    *blocked = new mpi_block_distribution
        (decomp,-1,globalsize);
mpi_object
    *grouping = new mpi_object( blocked );
```

The computation is a parallel map, for which for now we use the cross product kernel with null cross dimension:

```
%% unittest_kmeans.cxx
find_nearest_center = new mpi_outerproduct_kernel
( distances,grouping,NULL,group_calculation);
```

Eh,

```
%% unittest_kmeans.cxx
mpi_object
    *min_distance = new mpi_object( blocked );
```

Updating the centers takes two kernels. First we locally construct an array of size $N \times dk$ of coordinates, where only the κ -th coordinate is filled in, where κ is the number of the nearest center.

```

%% unittest_kmeans.cxx
mpi_distribution
    *k2blocked = new mpi_block_distribution
        (decomp,ncluster*dim,-1,globalsize);
mpi_object
    *masked_coordinates = new mpi_object( k2blocked );
%% unittest_kmeans.cxx
group_coordinates = new mpi_outerproduct_kernel
( coordinates,masked_coordinates,grouping,coordinate_masking );

```

This is yet another outer product kernel, with the grouping passed as context.

Next the new centers are computed by a reduction:

```

%% unittest_kmeans.cxx
mpi_distribution
    *klocal = new mpi_block_distribution(decomp,ncluster*dim,1,-1);
mpi_object
    *partial_sums = new mpi_object( klocal );
mpi_kernel
    *compute_new_centers1 = new mpi_kernel( masked_coordinates,partial_sums );
%% unittest_kmeans.cxx
compute_new_centers1->set_name("partial sum calculation");
compute_new_centers1->set_localexecutefn( &center_calculation_partial );
compute_new_centers1->last_dependency()->set_type_local();

```

See the theory of gather in section ??.

3 Discussion

References

- [1] 2010 ibm hpc challenge class ii submission.
- [2] T. F. Gonzalez. Clustering to minimize the maximum intercluster distance. *Theoretical Computer Science*, 38:293–306, 1985.
- [3] J. McClintock and A. Wirth. Efficient Parallel Algorithms for k-Center Clustering. *ArXiv e-prints*, April 2016.