

TACC Technical Report IMP-13

Sparse Operations in the Integrative Model for Parallelism

Victor Eijkhout*

February 22, 2016

This technical report is a preprint of a paper intended for publication in a journal or proceedings. Since changes may be made before publication, this preprint is made available with the understanding that anyone wanting to cite or reproduce it ascertains that no published version in journal or proceedings exists.

Permission to copy this report is granted for electronic viewing and single-copy printing. Permissible uses are research and browsing. Specifically prohibited are *sales* of any copy, whether electronic or hardcopy, for any purpose. Also prohibited is copying, excerpting or extensive quoting of any report in another work without the written permission of one of the report's authors.

The University of Texas at Austin and the Texas Advanced Computing Center make no warranty, express or implied, nor assume any legal liability or responsibility for the accuracy, completeness, or usefulness of any information, apparatus, product, or process disclosed.

* eijkhout@tacc.utexas.edu, Texas Advanced Computing Center, The University of Texas at Austin

Abstract

We consider the sparse matrix vector product.

The following IMP reports are available or under construction:

- IMP-00** The IMP Elevator Pitch
- IMP-01** IMP Distribution Theory
- IMP-02** The deep theory of the Integrative Model
- IMP-03** The type system of the Integrative Model
- IMP-04** Task execution in the Integrative Model
- IMP-05** Processors in the Integrative Model
- IMP-06** Definition of a ‘communication avoiding’ compiler in the Integrative Model
- IMP-07** Associative messaging in the Integrative Model (under construction)
- IMP-08** Resilience in the Integrative Model (under construction)
- IMP-09** Tree codes in the Integrative Model
- IMP-10** Thoughts on models for parallelism
- IMP-11** A gentle introduction to the Integrative Model for Parallelism
- IMP-12** K-means clustering in the Integrative Model
- IMP-13** Sparse Operations in the Integrative Model for Parallelism
- IMP-14** 1.5D All-pairs Methods in the Integrative Model for Parallelism (under construction)
- IMP-15** Collectives in the Integrative Model for Parallelism
- IMP-16** Processor-local code generation (under construction)
- IMP-17** The CG method in the Integrative Model for Parallelism (under construction)
- IMP-18** A tutorial introduction to IMP software (under construction)
- IMP-19** Report on NSF EAGER 1451204.
- IMP-20** A mathematical formalization of data parallel operations
- IMP-21** Adaptive mesh refinement (under construction)
- IMP-22** Implementing LULESH in IMP (under construction)

1 Introduction

In the formal treatment of Integrative Model for Parallelism (IMP) (see IMP-03) we introduced the ‘indirect’ function `Ind` mapping output indices to required input indices.

Let $x, y \in \text{Array}$ and let the function f compute each element of y from certain elements of x . We use a function σ (or σ_f to indicate its provenance explicitly) called the ‘signature function’ to determine the mapping from output indices to input sets of indices:

Datatype Signature: Signature of data parallel functions

$\text{Signature} \equiv N \rightarrow \text{Ind}$

As remarked, this function can be represented as a boolean sparse matrix, and for the case where f computes a matrix-vector product, that boolean matrix corresponds to the sparsity pattern of f .

2 Implementation

We define the sparse matrix-vector product as inheriting from a regular kernel, by setting the sparse matrix as both the pattern that will determine the beta distribution, and as context for the local function:

```
class mpi_spmvp_kernel : virtual public mpi_kernel {
public:
    mpi_spmvp_kernel( object *in, object *out, mpi_sparse_matrix *mat)
        : kernel(in, out), mpi_kernel(in, out) {
        set_name(fmt::format("sparse-mvp{}", get_out_object()->get_object_number()));
        dependency *d = last_dependency();
        d->set_index_pattern( mat );
        set_localexecutefn( &local_sparse_matrix_vector_multiply );
        localexecutectx = (void*)mat;
    };
    virtual void analyze_dependencies() override {
        mpi_kernel::analyze_dependencies();
        mpi_sparse_matrix *mat = (mpi_sparse_matrix*)localexecutectx;
        if (mat->get_trace())
            fmt::print("{}\n", get_out_object()->get_architecture()->mytid(), mat->as_string());
    };
};
```

The pattern is stored as as the local part of the beta distribution:

```
for (auto dom : gamma->get_domains()) {
    indexstruct *base = gamma->get_pstruct(dom);
    indexstruct *columns = pattern->all_columns_from(base);
```

```
beta_struct->set_pstruct( dom,columns );
}
```

which means that dependency analysis can be done as for an explicit beta:

```
%% imp_base.cxx
std::vector<message*> *distribution::messages_for_segment
( int mytid,int skipsel,int indexstruct *beta_block,indexstruct *halo_struct ) {
    index_int g,g_last; int P = this->global_ndomains();
    std::vector<message*> *messages = new std::vector<message*>;
    messages->reserve(P);
    indexstruct *buildup = new empty_indexstruct();
    int pstart = mytid;
    if (get_architecture()->has_random_sourcing())
        pstart += rand()%P;
    for (int ip=0; ip<P; ip++) {
        int p = (pstart+ip)%P; // start with self first
        if (!lives_on(p)) continue; // deal with masks
        std::shared_ptr<indexstruct> intersect { beta_block->intersect( processor_structure(p) ) };
        if (intersect->local_size()>0 && !(skipsel && p==mytid) && !buildup->contains(intersect.get())) {
            message *m = new message(p,mytid,intersect);
            if (get_architecture()->halo_has_local_address_space) m->relativize(halo_struct);
            messages->push_back( m );
            buildup = buildup->struct_union(intersect.get());
            if (buildup->contains(beta_block)) goto covered; //->equals(buildup) goto covered;
        }
    }
}
```

3 Remapping and caching

In the class definition of `mpi_spmvp_kernel` above we did not yet note the overloaded definition of `analyze_dependencies`: this includes a call to remap the matrix:

```
%% mpi_base.cxx
void mpi_sparse_matrix::remap( distribution *alpha,distribution *beta,int mytid ) {
    if (has_been_remapped) return;

    indexstruct *column_indices = beta->processor_structure(mytid);
    sparse_matrix::remap( 0,alpha->local_size(mytid)-1, column_indices );

    has_been_remapped = 1;
};
```

where MPI remapping is based on an internal routine:

```
void sparse_matrix::remap( index_int first, index_int last,indexstruct *column_indices ) {
    for (index_int irow=first; irow<=last; irow++) {
        indexstruct *oldrow = rows[irow], *newrow;
```

```
if (!rows[irow]->is_indexed())
    oldrow = rows[irow]->convert_to_indexed();

newrow = oldrow->make_clone();
for (index_int j=0; j<oldrow->local_size(); j++) {
    index_int col = oldrow->get_ith_element(j);
    index_int l;
    try { l = column_indices->find(col); }
    catch (std::string c) { fmt::printf("Error <<{}>>\n", c);
        newrow->set_ith_element(j, l);
    }
    rows[irow] = newrow;
}
};
```